

# Forecasting German Electricity Demand: A Comparison of Benchmark, SARIMAX, Feature-Based and Neural Forecasting Models

Mathujan Sivananthan - 24087340

## 1. Introduction

This report forecasts the electricity load in Germany with consideration of a wide variety of simple benchmarks to more and more complicated models, such as SARIMAX, feature-based tree models and neural networks. The research questions the models address are: (1) which models can meaningfully outperform a seasonal naive model; (2) how to prevent data leakage in features; and (3) why certain hyperparameter values are selected, with the differencing order being one example. In addition, we investigate if adding temperature covariates increases the accuracy of the forecast, examine the compromise between being easy to interpret and being more complex, and finally suggest a model that could be used in an operation. The main argument of this analysis is that it is very hard to beat a trivial seasonal naive benchmark fairly. However, on a real forecasting assignment on a fixed origin with a two-year-ahead forecast horizon, no complex model has any meaning when compared with it, mainly because of the structural change that occurred at the time of the COVID-19 pandemic.

## 2. Data and Preprocessing

The main dataset employed in this study is the Open Power System Data (OPSD) actual load data of Germany for 1 hour from January 2015 to October 2020. Data were cleaned, and missing data were removed, and the hourly load in Megawatts (MW) was converted to a weekly mean in Gigawatts (GW). This led to 301 weekly observations, of which 197 were assigned for training, and 104 observations were assigned for testing. The key issue here is the aggregation of the data to a weekly frequency, in order to remove the dominant daily and weekend cycles, and reveal a more pronounced annual seasonal signal. A strictly chronological train/test split was enforced to prevent any data leakage. In the case of the Long Short-Term Memory (LSTM) model, the *MinMaxScaler* was applied only to the training period of the hours, and then it was applied to the whole series. Furthermore, all the lag and rolling features were created from the shift operation of the tree-based models (e.g., `lag_1 = shift(1)`, `drop_first(52)` for `lag_52`), such that the features on week  $t$  only contain data until week  $t-1$ . It is important to note that the temperatures recorded during the test period represent future weather, and thus the temperatures used are, in this way, look-ahead data, rather than true operational ones, which makes the models used conditional forecasts.

### 3. Exploratory Analysis

This exploratory data analysis included visualising the raw series, performing seasonal and trend decomposition using Loess (STL), and examining the autocorrelation function (ACF) and partial autocorrelation function (PACF). The results of the honesty and critical tests on the level series were stationary, as indicated by the Dickey-Fuller test ( $p=0.001$ ) and the KPSS test ( $p=0.10$ ). The conclusion of this finding is interpreted as a strong deterministic annual seasonality as opposed to a stochastic unit root. This vital insight will give direction to the subsequent discussion of the suitable differencing orders for the SARIMA models.




Figure 1: Weekly average and Daily average German electricity load




Figure 2: STL decomposition of weekly load

#### 4. Forecasting Methods

Several model families were compared to standard benchmarks in the forecasting pipeline. A seasonal length of  $s=52$  was selected for the SARIMAX model due to the presence of an annual cycle as seen in the ACF plot with a spike at lag 52 and the STL decomposition. To eliminate this annual mean, an airline-style seasonal differencing ( $D=1$ ) was used. The initial tests of stationarity indicated that the non-seasonal differencing may not be required ( $d=0$ ), but the AIC grid search favoured a non-seasonal differencing of  $d=1$ . Adding  $d=1$  onto  $D=1$  for a stationary level series is a mild level of over-differencing, which plausibly increases the bias of the model's extrapolation (by +2.80 GW) but helps the in-sample information criteria. The models based on features (Random Forest and Gradient Boosting Machines) were organised into a recursive multi-step forecaster. Engineered features were used as inputs to the models, with shift-based lags. In the forecasting step, the models recursively injected their forecasts as lags, instead of real future load data. Finally, an LSTM neural network was implemented for shorter time horizons. The ability of LSTMs to model complex sequences of time without much manual feature engineering is well documented in the field of load forecasting (Hochreiter & Schmidhuber, 1997; Kong et al., 2019; Marino et al., 2016; Bouktif et al., 2018). This network was trained, however, for rolling forecasts over a week ahead and not set for a fixed two-year time horizon due to architectural requirements.

#### 5. Evaluation Design

The main evaluation task is at a fixed origin of 104 weeks. The models each provide a forecast at the end of the training period, and they each predict the forecast 2 years into the future. Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), Mean Absolute Scaled Error (MASE), and Bias were used for evaluating models. The MASE model serves as a reference benchmark model that is anchored with the in-sample model seasonal naive.

An important imbalance to highlight is that the LSTM model is predicting a week-ahead forecast that is rolled across the test window, re-anchoring on the latest actuals. This is a much simpler task than the fixed-origin 2-year forecast which the other models perform, so it cannot be fairly or directly compared to the others.

#### 6. Results

Table 1 shows the comparative accuracy on the 104-week test window. On the basis of RMSE, two models beat the baseline model from the seasonality (3.007 GW): first, the Random Forest (2.970 GW) and second, the LSTM (2.155 GW). The 1% advantage of the Random Forest, however, is statistically insignificant and lacking in any real substance. However, the model is based on lag\_52 for 72% of its feature importance,

which suggests that it was able to learn the seasonal naive benchmark better than any meaningful improvement upon it.

The LSTM's 28% deficit looks like a big figure but is very confounded by its weekly time frame. The seasonal naive benchmark is the best model on a fair like-for-like basis. Evidence is most robust in the Bias column, where all fixed-origin models' demand forecasts are highly positive (+1.69 to +2.80 GW). The seasonal naive model is very competitive because it straightforwardly copies the preceding season, and not the wrongly extrapolated past trends into an unobserved future.

| Model                             | RMSE  | MAE   | MASE  | Bias  |
|-----------------------------------|-------|-------|-------|-------|
| LSTM (week-ahead, rolled)*        | 2.155 | 1.657 | 1.238 | +0.53 |
| Feature model — Random Forest     | 2.97  | 2.212 | 1.604 | +1.69 |
| Seasonal naive (benchmark)        | 3.007 | 2.319 | 1.732 | +1.73 |
| Feature model — Gradient Boosting | 3.265 | 2.47  | 1.792 | +1.82 |
| SARIMAX + temperature             | 3.553 | 2.815 | 2.103 | +2.52 |
| SARIMA (2,1,6)(0,1,1,52)          | 3.762 | 3.08  | 2.301 | +2.80 |
| Mean                              | 4.397 | 3.789 | 2.831 | +0.48 |
| Naive                             | 4.459 | 3.783 | 2.827 | -0.88 |
| Drift                             | 5.118 | 4.34  | 3.243 | +1.01 |

Table 1: Weekly forecast accuracy on the 104-week test window (GW). Sorted by RMSE.

\* LSTM is a week-ahead forecast rolled across the test window (re-anchored on recent actuals), NOT a fixed-origin 2-year forecast — not directly comparable.




Figure 3: Seasonal Navie (Benchmark) forecasts




Figure 4: SARIMA forecast




Figure 5: Feature model forecast




Figure 6: LSTM week-ahead forecast, aggregated to weekly

## 7. Error Analysis

Looking at the errors in the forecast in chronological order reveals the primary problem: errors were not particularly large or volatile over 2019, but were very large in the spring of 2020. This separation is precisely at the time of the COVID-19 structural break. This is because every fixed-origin model over-forecasts (+1.7 to +2.8 GW bias) during the period, as it assumes a trend from before 2020 to after 2020, which it could not have anticipated in the macroeconomic disruption that occurred in 2020.

The LSTM can avoid this divergence because it keeps re-releasing to actuals. This was further evident in the probabilistic evaluations, where the confidence interval coverage

of the SARIMAX model dropped to about 51%, and the residual distributions were highly non-normal, with the Jarque-Bera p-value close to zero.

## **8. Discussion**

The use of temperature covariates with the SARIMAX model was not consistent. In-sample, it gave an excellent fit of the model (AIC reduced from 301 to 290) and solved the remaining Ljung-Box lag-12 p-value failure (from 0.005 to 0.21). But out of sample, it failed to outperform that seasonal naive benchmark, and that huge +2.52 GW bias remained. Only heating\_degree was significant ( $p=0.001$ ), while cooling\_degree was practically non-existent in Germany because the air conditioning load is negligible ( $p=0.88$ ) and temp\_mean was collinear. Furthermore, temperatures during a test period are technically "look-ahead" data, in that the weather that will occur during the test period cannot be forecast at the time the forecast is issued, and therefore are a conditional rather than an operational forecast. This study did not include holidays, which were not known from the calendar but instead represent an area of the predictive capability which was not captured.

Interpretability versus complexity is very simple. The SARIMAX model is very open to interpretation, including signed coefficients, principled confidence intervals, and formal diagnostics. The feature-based trees are in the centre and provide feature importances, but not directionality or inherent uncertainty. The LSTM is the highest complexity (requires extensive architecture, scaling, tuning and GPU compute) and the lowest interpretability (black box). Finally, there was no advantage of increased accuracy with increased complexity in this fixed-origin study.

The seasonal naive benchmark is the purest and most defensible recommendation for operational use for a 2-year-ahead planning horizon. It's played fairly, it doesn't need much upkeep, and it offers clear results. SARIMAX is a sensible improvement only when operators are very specific about wanting uncertainty intervals and/or explanations of weather impacts. However, on the other side of the equation, if the operational horizon is to be day/week ahead dispatch, a rolling SARIMAX or a rolling LSTM would be better, since their strength is with short-term re-anchoring.

## **9. Limitations and Future Work**

There are some constraints to this study. The geographic dimension of the temperature data did not average it over the whole country, but over one location (Berlin). Holiday features were not implemented and were a highly predictable source of variance. In addition, the 2020 structural change significantly affected the evaluation and tree models were structurally limited in their ability to predict outside of the training range, which is why they over-forecast 2020. Lastly, the uncertainty estimation of the methods

outside of the SARIMAX model was very weak, with the tree and LSTM intervals having unreliable coverage of 36–54%. Going forward, a rolling-origin cross-validation should be applied to examine robustness for each time horizon. Structural break can be addressed by COVID-19 dummy variables or intervention analysis. Model improvements might be made by introducing more probabilistic forecasting techniques, including the Fourier terms for SARIMA to better capture the long seasonality (52 seasons), and/or additional holiday covariates.

## **10. Conclusion**

When it comes to the hard but meaningful 2-year-ahead fixed-origin forecasting challenge, there is no simple model that does a better job of beating the trivial seasonal naive model. The COVID-19 pandemic creates a structural break, such that the test error overwhelmingly dominates, which leaves the advanced models unable to perform as expected because they extrapolate from trends before the pandemic. For this long-term task, the increased complexity of machine learning and neural architectures was not warranted. Hence, the seasonal naive model is suggested for 2 years.

## **References**

- Bouktif, S., Fiaz, A., Ouni, A. and Serhani, M.A. (2018). Optimal Deep Learning LSTM Model for Electric Load Forecasting... *Energies*, 11(7), 1636.
- Hersbach, H. et al. (2020). The ERA5 global reanalysis. *QJRMS*, 146, 1999–2049.
- Hochreiter, S. and Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8), 1735–1780.
- Hyndman, R.J. and Athanasopoulos, G. (2021). *Forecasting: Principles and Practice*, 3rd ed. OTexts.
- IEA (2020). *Power – The Covid-19 Crisis and Clean Energy Progress*. International Energy Agency.
- Kong, W., Dong, Z.Y., Hill, D.J., Luo, F. and Xu, Y. (2019). Short-Term Residential Load Forecasting Based on LSTM Recurrent Neural Network. *IEEE Transactions on Smart Grid*, 10(1), 841–851.
- Marino, D.L., Amarasinghe, K. and Manic, M. (2016). Building energy load forecasting using deep neural networks. *IECON 2016*, 7046–7051.
- Open Power System Data (2020). Time series data package (v2020-10-06).